# practical deep learning research and jobs

Alex Honchar University of Verona

#### Day 5 goals

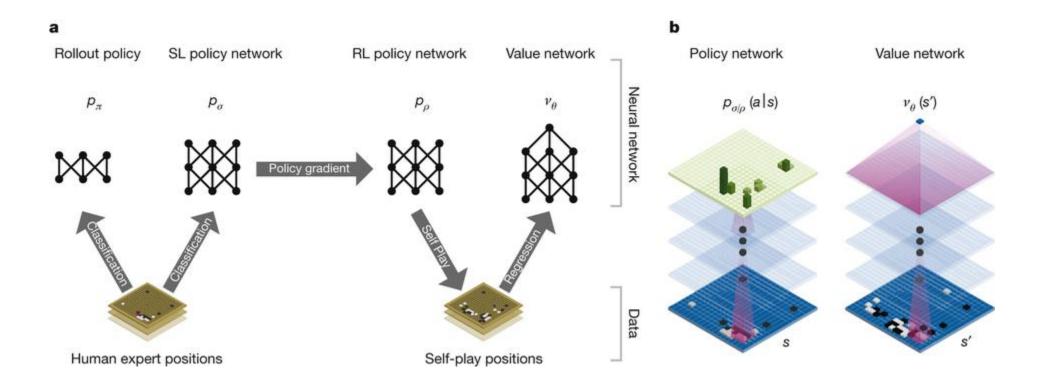
- You understand modern research trends in deep learning
- You understand how to get a job in machine learning area
- You can perform the full deep learning pipeline after making a final project

#### New technologies?

- Multilayer perceptron (1969)
- Convolutional neural networks (1980s)
- Recurrent neural networks (LSTM 1997)

# 2k16

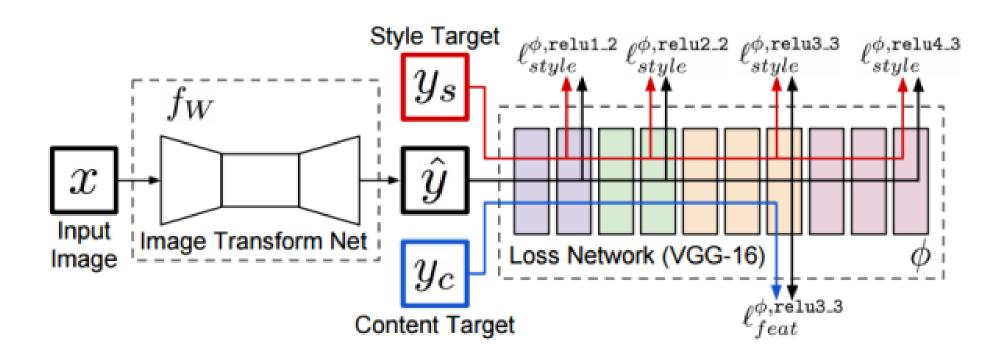
### 1. AlphaGo



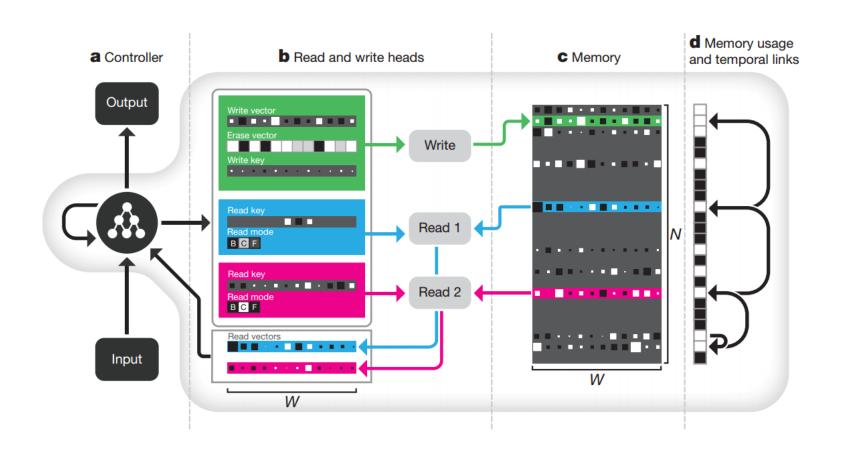
## 2. Image style transfer



#### 2. Image style transfer



#### 3. Differentiable neural computer

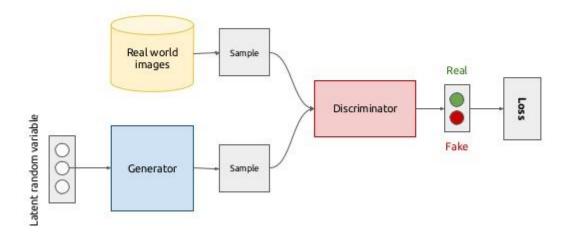


# 4. OpenAl



#### 5. Generative models

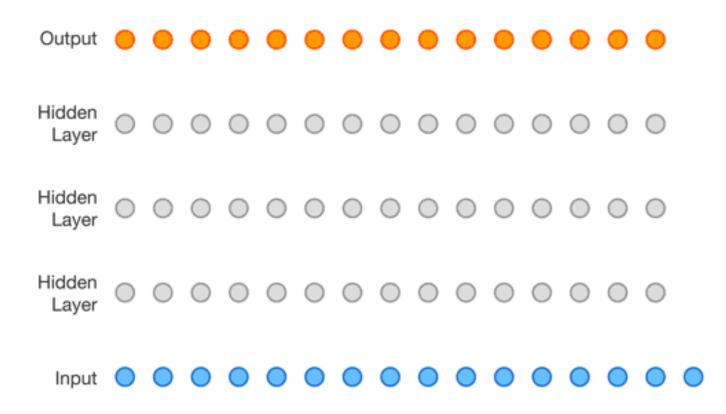
#### Generative adversarial networks (conceptual)



#### 6. Neural networks compression

Network	Top-1 Error	Top-5 Error	Parameters	Compress Rate
Baseline Caffemodel (BVLC)	42.78%	19.73%	240MB	1×
Fastfood-32-AD (Yang et al., 2014)	41.93%	-	131MB	$2\times$
Fastfood-16-AD (Yang et al., 2014)	42.90%	-	64MB	$3.7 \times$
Collins & Kohli (Collins & Kohli, 2014)	44.40%	-	61MB	$4 \times$
SVD (Denton et al., 2014)	44.02%	20.56%	47.6MB	$5 \times$
Pruning (Han et al., 2015)	42.77%	19.67%	27MB	$9 \times$
Pruning+Quantization	42.78%	19.70%	8.9MB	$27\times$
Pruning+Quantization+Huffman	42.78%	19.70%	6.9MB	<b>35</b> ×

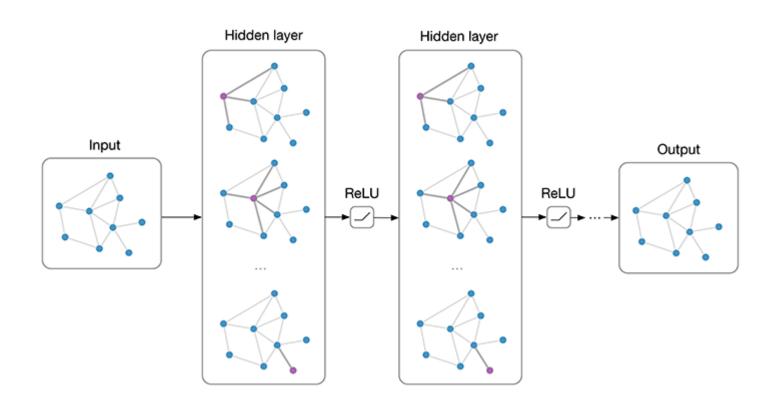
#### 7. Voice generation



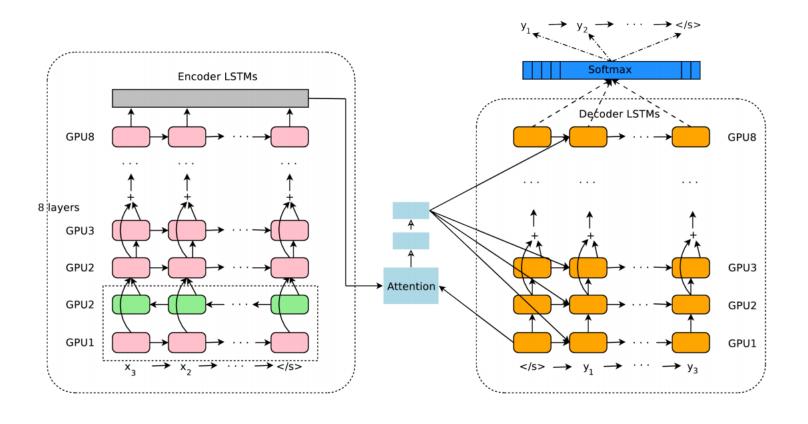
# 8. Donald Trump's win



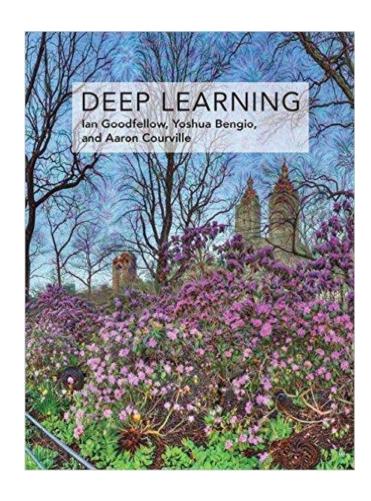
### 9. CNNs for graphs

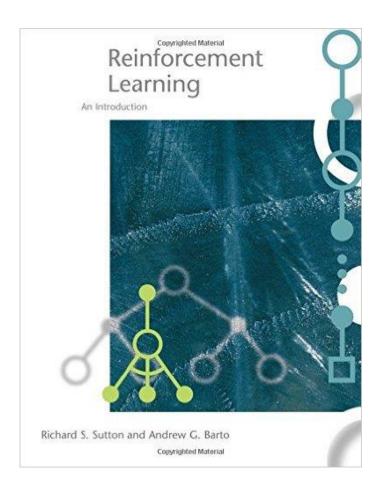


#### 10. End-to-end machine translation



#### + Good books





# jobs

#### Main professions

- Data Scientist
- Machine learning Engineer
  - Researcher

### Skills needed (computer science)

#### MUST:

- Programming skills (Python / C++ / Java / Matlab / Lua)
- Algorithms and data structures (discrete mathematics, graphs!)

#### IS A PLUS:

- Parallel computing
- Knowledge of GPU programming (CUDA)
- Backend development
- Data visualization skills

### Skills needed (mathematics)

#### MUST (2-3 years of bachelors in math/CS/statistics):

- Linear algebra (all about matrixes, vector spaces, linear operators)
- Calculus (limits, partial derivatives, series analysis)
- Probability theory (random variables, distributions, mean, variance)
- Statistics (central limit theorem, statistics tests, regression, Bayesian stats)

#### IS A PLUS:

- Numerical analysis (interpolations, numerical derivatives, numerical optimization)
- Tensor algebra
- Stochastic processes
- Multivariate statistics
- Functional analysis (for reading and understanding latest white papers)

#### 7 easy (no) steps to get a job

- 1. You must know and understand **basics** listed before
- 2. Complete one of the **machine learning** courses with certificate (Coursera Stanford course for example)
- 3. Make a submission on an easy Kaggle problem
- 4. Complete one or two courses / books in deep learning
- 5. Make a submission on intermediate / hard Kaggle problem
- 6. Implement one or two **white paper algorithms** from scratch using Theano / Tensorflow / Caffe and upload it to GitHub
- 7. Compile a resume and spam it on LinkedIn and head hunting sites

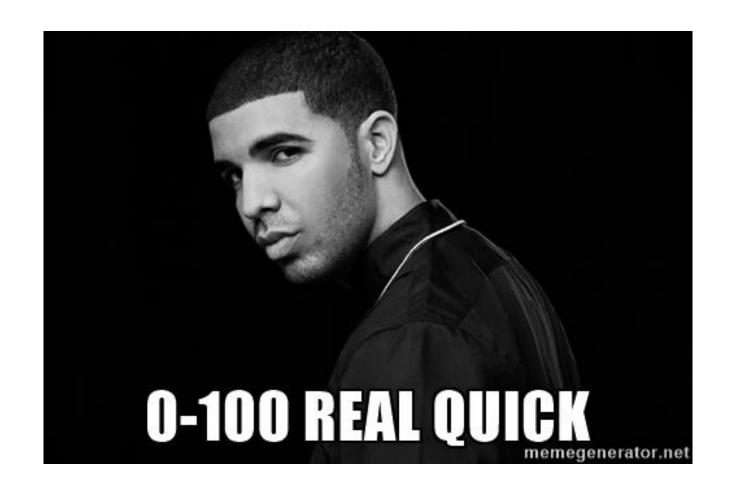
#### Steps MOOC mockup

- Python: <a href="https://www.codecademy.com/learn/python">https://www.codecademy.com/learn/python</a>
- Algorithms: <a href="https://www.coursera.org/learn/introduction-to-algorithms">https://www.coursera.org/learn/introduction-to-algorithms</a>
- Algebra: <a href="https://www.mooc-list.com/course/ut501x-linear-algebra-foundations-frontiers-edx">https://www.mooc-list.com/course/ut501x-linear-algebra-foundations-frontiers-edx</a>
- Calculus: <a href="https://www.coursera.org/learn/calculus1">https://www.coursera.org/learn/calculus1</a>
- Statistics: <a href="https://www.coursera.org/specializations/statistics">https://www.coursera.org/specializations/statistics</a>
- Machine learning: <a href="https://www.coursera.org/learn/machine-learning">https://www.coursera.org/learn/machine-learning</a>
- Kaggle 1: <a href="https://www.kaggle.com/c/titanic">https://www.kaggle.com/c/titanic</a>
- Deep Learning 1 (computer vision): <a href="http://cs231n.stanford.edu/">http://cs231n.stanford.edu/</a>
- Deep Learning 2 (natural language processing): <a href="http://cs224d.stanford.edu/">http://cs224d.stanford.edu/</a>
- Kaggle 2: any from "featured" section
- White paper 1: <a href="http://arxiv.org/pdf/1603.08155v1.pdf">http://arxiv.org/pdf/1603.08155v1.pdf</a> (image style transfer project)
- White paper 2: <a href="https://arxiv.org/pdf/1609.08359v1.pdf">https://arxiv.org/pdf/1609.08359v1.pdf</a> (emoji project)

#### Next steps

- Read white papers. Every week, better every day. Notice the trends, use them. (arxiv, gitxiv, Twitter)
- Implement models (at least their parts) from scratch for educational purpose (implement custom convolutional layer for Keras)
- 3. Probably dive a bit more into theory
- 4. Practice with data: do all the "dirty job"
- 5. Publish papers / commit to Github

#### Self-study materials from zero



Find your PDF here: <a href="https://github.com/Rachnog/education">https://github.com/Rachnog/education</a>

## Project proposals

- Titanic survivors prediction
  - House price prediction
- Fashion image classification
  - Food image classification
  - Arrhythmia ECG detection
- Plants images classification
- Adult content recognition
  - Stock price prediction

- Face identification
- Motivational quotes generation
  - Chat bots
- Speech recognition
- Machine translation
- Stealing ML models from API
  - Your ideas

thank you for attention!